

ARCHIVED COMPARISON, REPRODUCIBLE WITHOUT AN API KEY

Untreated fixtures pass 0/12. Baseline is one direct structured-output call per case (claude-haiku-4-5, 3 sweeps). Agent is the iterative tool-using workflow (claude-haiku-4-5, 3 sweeps). Both receive the identical task contract and fixtures, the deliberate delta is tools plus iteration.

HEADLINE METRICS

METRIC	SIMPLE BASELINE	AGENT SOLUTION	CHANGE
Pass rate per sweep (all three gates)	4/12, 7/12, 5/12	11/12, 12/12, 11/12	see per-case table
Worst sweep (reliability floor)	4/12	11/12	the deployment-honest number
Cases passing every sweep	4	11	verification converts flaky into stable
Machine time per case, zero human minutes in both arms	7.4s	110.1s	unattended either way
Cost per sweep of 12, client-side estimate	\$0.093	\$0.697	quality is the trade

PER-CASE OUTCOMES (PASSES ACROSS SWEEPS)

CASE	UNTREATED AXE	BASELINE	AGENT
01-missing-labels	4	3/3	3/3
02-low-contrast	9	0/3	3/3
03-div-buttons	0	3/3	3/3
04-missing-alt	4	3/3	3/3
05-heading-order	2	1/3	3/3
06-focus-indicators	0	3/3	3/3
07-modal	0	0/3	1/3
08-form-errors	3	0/3	3/3
09-custom-dropdown	0	0/3	3/3
10-landmarks	5	0/3	3/3
11-lang-title	2	1/3	3/3
12-tabindex-order	5	2/3	3/3

CONSISTENCY DETAIL

- Baseline never passes: 02-low-contrast, 07-modal, 08-form-errors, 09-custom-dropdown, 10-landmarks
- Baseline flaky: 05-heading-order, 11-lang-title, 12-tabindex-order
- Agent never passes: none
- Agent flaky: 07-modal

READING THE TABLE

- Cases 03, 06, 07 and 09 carry near-zero axe counts while seeding real keyboard failures, they measure what a linter cannot see.
- Every number above is recomputed by the harness from the stored working copies, run `npm run verify` to reproduce it locally without any API key.
- Full detail per run lives next to each run in `results.md` and `results.json`, per-case evidence in `score.json` and trajectory files.